

M2 Progress Report – Faithful RAG: Evaluating Citation Accuracy and Retrieval Robustness in Domain-Specific Scientific Literature

Andrea Trugenberger | 357615 | andrea.trugenberger@epfl.ch

Elie Bruno | 355932 | elie.bruno@epfl.ch

Faruk Zahiragić | 415360 | faruk.zahiragic@epfl.ch

Yusif Askari | 413862 | yusif.askari@epfl.ch

Team CiteRight

1 Progress Summary

All four proposal components run end-to-end on a 37-question gold benchmark of intellectual-property and innovation-policy papers. The retrieval ablation grew from four to sixteen configurations with a new corpus-level SOTA (§2); the faithfulness verifier exposes a 55–61 pp gap between NLI and LLM-judge labels (§3); RAGAS pits chunked RAG against a 1M-token long-context baseline and adds an answer-LLM cost dimension (§4); the Corrective RAG router lifts retrieval correctness from 60% to 80% on a mixed-quality 10-question slice by rescuing both ambiguous retrievals (§5).

2 Retrieval Ablation

The 37-question gold set (5 categories incl. multi-hop, byte-validated character spans, 6 adversarial controls q900–q905 that break the single-class Cohen- κ paradox (Cohen, 1960) from the first IAA round) drives a sixteen-configuration ablation. Configurations cross four embedder families (OpenAI text-embedding-3-small; BGE-M3 (Chen et al., 2024); E5-large (Wang et al., 2024); ColBERTv2 (Santhanam et al., 2022)), two chunk granularities (coarse $\sim$ 2000 chars, fine $\sim$ 300 chars), and an optional ZeroEntropy cross-encoder reranker, sharing one HybridRetriever code path that reports hit@ k , MRR, and nDCG.

config (coarse_rerank)	hit@5	hit@10	MRR
OpenAI text-emb-3-small	0.946	0.946	0.842
BGE-M3	0.946	0.946	0.842
E5-large	0.973	0.973	0.878
ColBERTv2	0.892	0.919	0.786

Table 1: Best per embedder; full grid in retrieval_eval/.

E5-large with the reranker wins every paper-level metric, beating the closed-source OpenAI baseline by +8.6 pp MRR; ColBERTv2 trails by 17 pp MRR. Multi-hop hit@10 falls from 1.000 on coarse to 0.667–0.833 on fine: the chunker is the dominant bottleneck.

3 Faithfulness Verification

The verifier extends the proposal’s NLI pipeline (DeBERTa-v3-large-mnli, He et al., 2021) with an LLM-judge head. On 24 answerable gold questions answered closed-book by three generators (GPT-4o-mini, Claude-3.5-haiku, DeepSeek-chat), NLI marks 18/24 (0.750) faithful at every provider while the LLM judge marks 0.138/0.173/0.198 — a 55–61 pp gap matching a known NLI failure mode (Min et al., 2023; Gao et al., 2023): NLI alone fails as a faithfulness gate.

4 End-to-End RAG vs Long-Context

Two pipelines run through the four RAGAS metrics (Es et al., 2024) on an $n=8$ stratified subset: chunked RAG (§2 retriever, GPT-4o-mini) and a long-context baseline that loads every cited document.md into a 1M-token Gemini 2.5 Flash window, judged by GPT-4o-mini. LC wins every metric but gaps are small (RAG/LC): faithfulness 0.899/0.933, answer_relevancy 0.713/0.789, context_precision 1.000/1.000, context_recall 0.917/1.000 — the open recall gap replicates the chunker bottleneck from §2. OpenRouter cost captures show LC at $\sim 10\times$ the answer-LLM spend (\$0.00383 vs \$0.00038/query), reframing the trade-off as accuracy-per-dollar.

5 Corrective RAG and M3 Plan

The CRAG router (Yan et al., 2024) scores top- k retrievals as CORRECT/AMBIGUOUS/INCORRECT and re-queries on AMBIGUOUS. On a mixed 10-question slice ($\tau_h=0.813$), CRAG lifts correct-retrieval from 60% to 80% (+20 pp), rescuing both ambiguous queries in one refinement round. A 9-cell (τ_h, τ_l) sweep hits the same 0.80 rate but spans 0.04–0.40 s/query, with ($\tau_h=0.7, \tau_l=0.4$) latency-optimal. M3 scales to 37 questions with RAGAS faithfulness scoring and an INCORRECT-branch fallback.

References

- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. 2024. [BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation](#).
- Jacob Cohen. 1960. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46.
- Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. 2024. [RAGAS: Automated evaluation of retrieval augmented generation](#).
- Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. 2023. [Enabling large language models to generate text with citations](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*.
- Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2021. [DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing](#).
- Sewon Min, Kalpesh Krishna, Xinxi Lyu, Mike Lewis, Wen tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. [FACTScore: Fine-grained atomic evaluation of factual precision in long form text generation](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*.
- Keshav Santhanam, Omar Khattab, Jon Saad-Falcon, Christopher Potts, and Matei Zaharia. 2022. [ColBERTv2: Effective and efficient retrieval via lightweight late interaction](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics*.
- Liang Wang, Nan Yang, Xiaolong Huang, Binxing Jiao, Linjun Yang, Daxin Jiang, Rangan Majumder, and Furu Wei. 2024. [Text embeddings by weakly-supervised contrastive pre-training](#).
- Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. 2024. [Corrective retrieval augmented generation](#).