

Progress Report: Faithful RAG – Evaluating Citation Accuracy and Retrieval Robustness in Domain-Specific Scientific Literature

Elie Bruno | 355932 | elie.bruno@epfl.ch
 Andrea Trugenberger | SCIPER | andrea.trugenberger@epfl.ch
 Faruk Zahiragić | SCIPER | faruk.zahiragic@epfl.ch
 Yusif | SCIPER | yusif@epfl.ch
 Team Faithful RAG

1 Progress Summary

Three of the four proposal components run end-to-end on a 37-question gold benchmark of intellectual-property and innovation-policy papers. The retrieval ablation grew from four to sixteen configurations (new corpus-level SOTA); the faithfulness verifier exposes a 55–61 pp gap between NLI and LLM-judge labels; the RAGAS pipeline pits chunked RAG against a 1M-token long-context baseline; the Corrective RAG hook is wired and awaits threshold ablations in M3.

2 Gold Benchmark

Each team member contributed ~10 question-passage pairs across policy impact, methodology, comparison, factual extraction, and multi-hop reasoning. A pair encodes an atomic claim, the source `paper_id`, character-span offsets on a line-feed-normalised `document.md`, and a verbatim quote. A validator byte-checks every span against its quote, and CI rejects fabricated concatenations. Inter-annotator agreement uses five pairs per annotator labelled *supports/contradicts/unrelated*; six adversarial controls (q900–q905) populate the rejection cells and break the single-class Cohen- κ paradox (Cohen, 1960) that the first IAA round triggered.

3 Retrieval Ablation

The ablation crosses four embedder families (OpenAI `text-embedding-3-small`; BGE-M3 (Chen et al., 2024); E5-large (Wang et al., 2024); ColBERTv2 (Santhanam et al., 2022)), two chunk granularities (coarse ~2000 chars, fine ~300 chars), and an optional ZeroEntropy cross-encoder reranker. Sixteen configurations share one HybridRetriever code path and report `hit@k`, MRR, and nDCG over $n=37$ paper-level gold pairs.

E5-large with the reranker wins every paper-level metric and beats OpenAI even without it (+10.8 pp `hit@10`, +8.6 pp MRR): a 335M-

config (coarse_rerank)	hit@5	hit@10	MRR
OpenAI <code>text-emb-3-small</code>	0.946	0.946	0.842
BGE-M3	0.946	0.946	0.842
E5-large	0.973	0.973	0.878
ColBERTv2	0.892	0.919	0.786

Table 1: Best configuration per embedder family. Full 16-config grid lives in the repo `retrieval_eval/` results.

parameter open-source encoder edges out the closed-source baseline. ColBERTv2 trails E5-large by 17 pp MRR despite a more expressive per-token similarity; its MS-MARCO pretraining fails to transfer to scholarly text. ColBERTv2 is the only family that prefers fine chunks. Multi-hop questions saturate `hit@10` at 1.000 on coarse and fall to 0.667–0.833 on fine across every family, locating the chunker as the dominant bottleneck.

4 Faithfulness Verification

The verifier extends the proposal’s NLI pipeline (DeBERTa-v3-large-mnli, He et al., 2021) with an LLM-judge head. Three generators (GPT-4o-mini, Claude-3.5-haiku, DeepSeek-chat) answer the 24 answerable gold questions closed-book; both verifiers label each (claim, gold-passages) pair.

generator	NLI	LLM-judge	gap
gpt-4o-mini	0.750	0.138	+0.612
claude-3.5-haiku	0.750	0.173	+0.577
deepseek-chat	0.750	0.198	+0.552

Table 2: Fraction labelled *faithful* by each verifier ($n=24$).

NLI labels three-quarters of hallucinated closed-book answers as faithful at every provider; the LLM judge flags them. The 55–61 pp gap is stable across generators and replicates a known NLI failure mode (Min et al., 2023; Gao et al., 2023). NLI alone fails as a faithfulness gate on this benchmark.

5 End-to-End RAG vs Long-Context

Two pipelines run through the four RAGAS metrics (Es et al., 2024): chunked RAG (the §3 retriever, GPT-4o-mini) and a long-context baseline

that loads every cited `document.md` into a 1M-token Gemini 2.5 Flash window.

metric	RAG	long-context	Δ
faithfulness	0.806	0.975	+0.169
answer_relevancy	0.662	0.799	+0.138
context_precision	0.979	1.000	+0.021
context_recall	0.792	1.000	+0.208

Table 3: RAGAS metrics, $n=8$ stratified subset.

Long-context wins every metric. The widest gaps fall on `context_recall` (+20.8 pp) and `faithfulness` (+16.9 pp): chunked RAG hides ~ 17 pp of available factual support behind chunk boundaries. Two metrics point at one root cause and reinforce the chunker bottleneck from Table 1.

6 Limitations and M3 Plan

Three scope ceilings remain: RAGAS at $n=8$, partial chunk-level coverage (16–20 of 37 queries), and a closed-book versus retrieved-context split between the two faithfulness verifiers. M3 delivers four fixes. (i) A no-gap sliding-window chunker: the multi-hop fine-chunk failures, partial chunk coverage, and the +17 pp long-context advantage share one root cause. (ii) Domain-adaptation of ColBERTv2 on synthetic contrastive pairs from the gold set, sized to close the 17 pp MRR gap. (iii) The full 37-question RAGAS sweep across all 16 retriever configurations. (iv) Activation of the wired CRAG (Yan et al., 2024) scorer as a 17th configuration aimed at the multi-hop `fine` failures isolated in §3.

References

- Jianlv Chen, Shitao Xiao, Peitian Zhang, Kun Luo, Defu Lian, and Zheng Liu. 2024. [BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation](#).
- Jacob Cohen. 1960. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46.
- Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. 2024. [RAGAS: Automated evaluation of retrieval augmented generation](#).
- Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. 2023. [Enabling large language models to generate text with citations](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*.
- Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2021. [DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing](#).
- Sewon Min, Kalpesh Krishna, Xinxi Lyu, Mike Lewis, Wen tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. [FactScore: Fine-grained atomic evaluation of factual precision in long form text generation](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*.
- Keshav Santhanam, Omar Khattab, Jon Saad-Falcon, Christopher Potts, and Matei Zaharia. 2022. [ColBERTv2: Effective and efficient retrieval via lightweight late interaction](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics*.
- Liang Wang, Nan Yang, Xiaolong Huang, Binxing Jiao, Linjun Yang, Daxin Jiang, Rangan Majumder, and Furu Wei. 2024. [Text embeddings by weakly-supervised contrastive pre-training](#).
- Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. 2024. [Corrective retrieval augmented generation](#).